

Kubernetes 节点弹性伸缩开源组件 Karpenter 实验

1. 配置环境变量

```
export CLUSTER_NAME="karpenter-clu"

export AWS_REGION="us-west-2"

export AWS_ACCOUNT_ID="$(aws sts get-caller-identity --query Account --output text)"
```

2. 创建 Cluster 的配置文件

```
cat << EOF > cluster.yaml
apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: ${CLUSTER_NAME}
  region: ${AWS_REGION}
  version: "1.21"
  tags:
    karpenter.sh/discovery: ${CLUSTER_NAME}

iam:
  withOIDC: true

managedNodeGroups:
  - name: ng-1
    privateNetworking: true
    instanceType: m5.large
    desiredCapacity: 3
EOF
```

3. 安装 eksctl 和 kubelet

```
# eksctl
curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp

sudo mv -v /tmp/eksctl /usr/local/bin

# kubelet
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubect1"

curl -LO https://dl.k8s.io/release/v1.21.0/bin/linux/amd64/kubect1

sudo install -o root -g root -m 0755 kubect1 /usr/local/bin/kubect1

sudo curl --silent --location -o /usr/local/bin/kubect1 \
  https://amazon-eks.s3.us-west-2.amazonaws.com/1.19.6/2021-01-05/bin/linux/amd64/
kubect1

sudo chmod +x /usr/local/bin/kubect1
```

4. 部署 cluster

```
eksctl create cluster -f cluster.yaml

export CLUSTER_ENDPOINT="$(aws eks describe-cluster --name ${CLUSTER_NAME} --query
"cluster.endpoint" --output text)"

echo $CLUSTER_ENDPOINT

[https://4F71450F00F2D92463DDE6BDF0A64B4C.gr7.us-west-2.eks.amazonaws.com](https://
4F71450F00F2D92463DDE6BDF0A64B4C.gr7.us-west-2.eks.amazonaws.com)
```

5. 配置 karpenter 用户角色

```
TEMPOUT=$(mktemp)
```

```
curl -fsSL https://karpenter.sh/v0.6.3/getting-started/cloudformation.yaml > $TEMPOUT \
&& aws cloudformation deploy \
  --stack-name "Karpenter-${CLUSTER_NAME}" \
  --template-file "${TEMPOUT}" \
  --capabilities CAPABILITY_NAMED_IAM \
  --parameter-overrides "ClusterName=${CLUSTER_NAME}"
```

```
eksctl-karpenter-addon-iamserviceaccount-kub-Role1-1G9PX8G86AXUL
```

身份提供商: arn:aws:iam::229474920661:oidc-provider/oidc.eks.us-west-2.amazonaws.com/id/
7D0BD76F276A50872F4C306BEA8FDFB3

```
eksctl-karpenter-cluster-ServiceRole-1K4XCJW9OGGI8
```

AWS 服务: eks

```
eksctl-karpenter-nodegroup-ng-1-NodeInstanceRole-1O8B30SYRDXT3
```

AWS 服务: ec2

```
KarpenterNodeRole-karpenter
```

AWS 服务: ec2

6. 配置 iamidentitymapping 和 iamserviceaccount

```
eksctl get iamidentitymapping --cluster ${CLUSTER_NAME}
ARN
USERNAME                                GROUPS                                ACCOUNT
arn:aws:iam::229474920661:role/eksctl-karpenter-clu-nodegroup-ng-NodeInstanceRole-03NAG1C0ZA4Q system:node:{{
EC2PrivateDNSName}}                  system:bootstrappers,system:nodes

eksctl get iamserviceaccount --cluster ${CLUSTER_NAME}
NAMESPACE      NAME      ROLE ARN
kube-system    aws-node  arn:aws:iam::229474920661:role/
eksctl-karpenter-clu-addon-iamserviceaccount-Role1-9AA01F3X1KDS
```

```
# 关联role
eksctl create iamidentitymapping \
  --username system:node:{{EC2PrivateDNSName}} \
  --cluster "${CLUSTER_NAME}" \
  --arn "arn:aws:iam::${AWS_ACCOUNT_ID}:role/KarpenterNodeRole-${CLUSTER_NAME}" \
  --group system:bootstrappers \
  --group system:nodes

#
eksctl utils associate-iam-oidc-provider --cluster ${CLUSTER_NAME} --approve

#
eksctl create iamserviceaccount \
  --cluster "${CLUSTER_NAME}" --name karpenter --namespace karpenter \
  --role-name "${CLUSTER_NAME}-karpenter" \
  --attach-policy-arn "arn:aws:iam::${AWS_ACCOUNT_ID}:policy/KarpenterControllerPolicy-${CLUSTER_NAME}" \
  --role-only \
  --approve
```

```
eksctl get iamidentitymapping --cluster ${CLUSTER_NAME}

ARN                               USERNAME                               GROUPS                               ACCOUNT
arn:aws:iam::229474920661:role/KarpenterNodeRole-karpenter-clu system:node:{{
EC2PrivateDNSName}}               system:bootstrappers,system:nodes
arn:aws:iam::229474920661:role/eksctl-karpenter-clu-nodegroup-ng-NodeInstanceRole-03NAG1C0ZA4Q system:node:{{
EC2PrivateDNSName}}               system:bootstrappers,system:nodes

#####
eksctl get iamserviceaccount --cluster ${CLUSTER_NAME}
NAMESPACE   NAME      ROLE ARN
karpenter   karpenter arn:aws:iam::229474920661:role/karpenter-clu-karpenter
kube-system aws-node  arn:aws:iam::229474920661:role/
eksctl-karpenter-clu-addon-iamserviceaccount-Role1-9AA01E3XLKDS
```

```
export KARPENTER_IAM_ROLE_ARN="arn:aws:iam::${AWS_ACCOUNT_ID}:role/${CLUSTER_NAME}-karpenter"
```

```
echo $KARPENTER_IAM_ROLE_ARN
arn:aws:iam::229474920661:role/karpenter-clu-karpenter
```

7. 安装 helm 和 karpenter

```
# 安装helm
curl -sSL https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 | bash

# 安装
helm repo add karpenter https://charts.karpenter.sh/
helm repo update

helm upgrade --install --namespace karpenter --create-namespace \
  karpenter karpenter/karpenter \
  --version v0.10.1 \
  --set serviceAccount.annotations."eks\.amazonaws\.com/role-arn"=$
{KARPENTER_IAM_ROLE_ARN} \
  --set clusterName=${CLUSTER_NAME} \
  --set clusterEndpoint=${CLUSTER_ENDPOINT} \
  --set aws.defaultInstanceProfile=KarpenterNodeInstanceProfile-${CLUSTER_NAME} \
  --set logLevel=debug \
  --wait

kubectl get deployment -n karpenter

kubectl get pods -n karpenter
```

8. 配置 Provisioner

```
cat << EOF > provisioner.yaml
apiVersion: karpenter.sh/v1alpha5
kind: Provisioner
metadata:
  name: cpu
spec:
  requirements:
    - key: karpenter.sh/capacity-type
      operator: In
      values: ["spot"]
    - key: node.kubernetes.io/instance-type
      operator: In
      values: ["m5.xlarge", "m5.2xlarge"]
  limits:
    resources:
      cpu: 1000
  provider:
    subnetSelector:
      karpenter.sh/discovery: karpenter-clu
    securityGroupSelector:
      kubernetes.io/cluster/karpenter-clu: owned
    ttlSecondsAfterEmpty: 30
EOF
```

```
kubectl get provisioners
```

```
kubectl describe provisioners
```

9. 使用 karpenter 扩容

```
# 扩容 deployment
kubectl scale deployment eks-sample-linux-deployment --replicas 200 -n eks-sample-app

# 缩容
kubectl scale deployment eks-sample-linux-deployment --replicas 0 -n eks-sample-app

# 查看 karpenter 日志
kubectl logs -f -n karpenter -l app.kubernetes.io/name=karpenter -c controller
```